

MACHINE LEARNING IN ROBOTICS

Assignment 2

Stefan Röhl

Technische Universität München, Arcisstraße 21, Munich, Germany

Email: stefan.roehrl@tum.de

I. EXERCISE

Learned parameters for GMM after
Expectation-Maximization.

Covariance Matrices:

$$\Sigma_1 = \begin{pmatrix} 3.946e-04 & 2.169e-04 \\ 2.169e-04 & 1.277e-04 \end{pmatrix}$$

$$\Sigma_2 = \begin{pmatrix} 7.436e-04 & -5.914e-04 \\ -5.914e-04 & 6.098e-04 \end{pmatrix}$$

$$\Sigma_3 = \begin{pmatrix} 1.083e-03 & -4.244e-04 \\ -4.244e-04 & 2.431e-04 \end{pmatrix}$$

$$\Sigma_4 = \begin{pmatrix} 1.748e-04 & 2.616e-04 \\ 2.616e-04 & 3.976e-04 \end{pmatrix}$$

Means:

$$\mu_1 = \begin{pmatrix} -1.470e-02 \\ -7.963e-02 \end{pmatrix}$$

$$\mu_2 = \begin{pmatrix} -1.937e-02 \\ -1.664e-02 \end{pmatrix}$$

$$\mu_3 = \begin{pmatrix} 2.619e-02 \\ 6.173e-02 \end{pmatrix}$$

$$\mu_4 = \begin{pmatrix} -4.319e-02 \\ 4.459e-02 \end{pmatrix}$$

Priors:

$$\pi_1 = (2.012e-01)$$

$$\pi_1 = (2.971e-01)$$

$$\pi_1 = (2.617e-01)$$

$$\pi_1 = (2.400e-01)$$

The log-likelihood before the EM is: 1.353e+03

The log-likelihood after the EM is: 1.448e+03

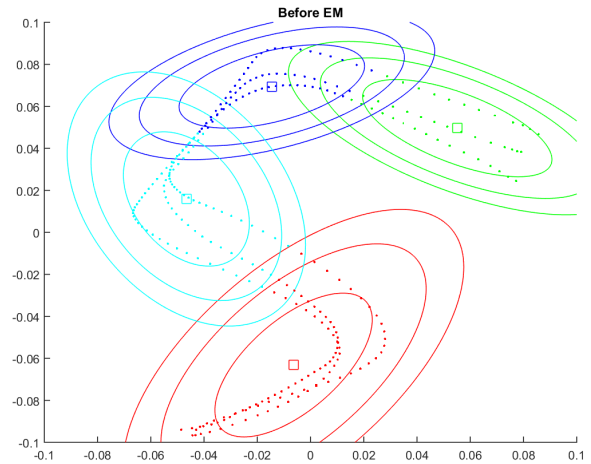


Figure 1. pdf and clusters before EM

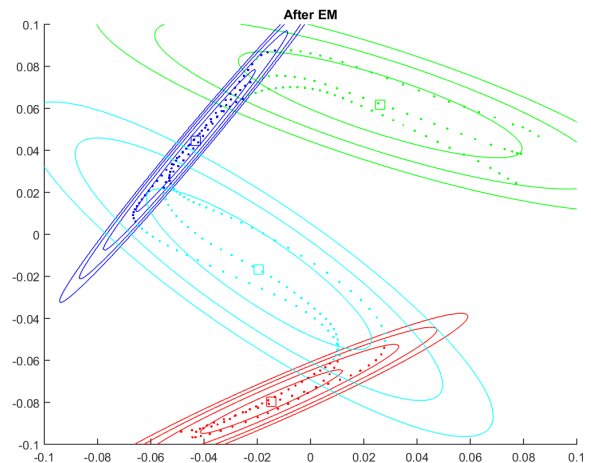


Figure 2. pdf and clusters after EM

II. EXERCISE

The following list contains the log-likelihood and the labels of the Train Set. Label 1 stands for class *train* and 0 stands for class *test*.

$$[\log - \text{likelihood}, \text{label}] =$$

-2.312	1.000
-2.312	1.000
-14.620	1.000
-24.359	1.000
-46.196	1.000
-46.196	1.000
-70.710	1.000
-87.584	1.000
-68.024	1.000
-67.401	1.000
-89.607	1.000
-92.499	1.000
-126.337	0.000
-115.001	1.000
-70.706	1.000
-94.465	1.000
-115.794	1.000
-95.986	1.000
-117.905	1.000
-81.088	1.000
-89.167	1.000
-89.331	1.000
-50.709	1.000
-80.423	1.000
-57.664	1.000
-40.420	1.000
-56.434	1.000
-45.929	1.000
-45.789	1.000
-39.312	1.000
-39.312	1.000
-39.172	1.000
-38.899	1.000
-38.348	1.000
-38.348	1.000
-38.070	1.000
-37.936	1.000
-37.803	1.000
-37.803	1.000
-37.669	1.000
-37.193	1.000
-36.850	1.000
-36.583	1.000
-36.449	1.000
-35.973	1.000
-35.630	1.000
-35.630	1.000
-35.287	1.000
-35.287	1.000
-34.944	1.000
-34.601	1.000
-34.258	1.000
-34.258	1.000
-24.866	1.000
-24.866	1.000
-24.866	1.000
-2.312	1.000
-2.312	1.000
-2.312	1.000
-2.312	1.000

This list contains the log-likelihood of the Test Set. Where again 1 stands for class *train* and 0 stands for class *test*. As the result shows all test have a log-likelihood value over -120 so all are classified as *train*.

$$[\log - \text{likelihood}, \text{label}] =$$
[illegible]

III. EXERCISE

• Policy Iteration:

Reward Matrix

$$R = \begin{pmatrix} 0.000 & 0.000 & 0.000 & 0.000 \\ 0.000 & 1.000 & -10.000 & -5.000 \\ 0.000 & 0.000 & -10.000 & -5.000 \\ 0.000 & 0.000 & 0.000 & 0.000 \\ -10.000 & -5.000 & 0.000 & 1.000 \\ 0.000 & 0.000 & 0.000 & 0.000 \\ 0.000 & 0.000 & 0.000 & 0.000 \\ -10.000 & 5.000 & 0.000 & 0.000 \\ -10.000 & -5.000 & 0.000 & 0.000 \\ 0.000 & 0.000 & 0.000 & 0.000 \\ 0.000 & 0.000 & 0.000 & 0.000 \\ -10.000 & 5.000 & 0.000 & 0.000 \\ 0.000 & 0.000 & 0.000 & 0.000 \\ 0.000 & 0.000 & -10.000 & 5.000 \\ 0.000 & 0.000 & -10.000 & 5.000 \\ 0.000 & 0.000 & 0.000 & 0.000 \end{pmatrix}$$

Used Value of γ : The used discount factor is held similar to the exercise 6 close to 1. Values from 0.99 to 0.999 were used.

Iterations required: Depending on the random start values the algorithm need 5-9 Iterations till convergence.

To have a better surveillance over the states which were passed I modified the walkshow function in a way that the state number is printed under the image.

Results of *WalkPolicyIteration(s)*:




Figure 3. Results from starting state $s = 8$



Figure 4. Results from starting state $s = 10$




Figure 5. Results from starting state $s = 3$

• Q-Learning:

Used ϵ : As the task says the probability ϵ should be small values between 0.1 and 0.2 are used. These values are static and don't change during the procedure. I am not sure if i found the ideal policy, but i think to improve the algorithm one should at first keep ϵ small to achieve fast convergence. After some steps when the algorithm has found a good policy one could increase ϵ to let the robot try more random actions. This should help to get out of local convergence to global convergence. I think the optimal policy could be found more often with this adaption, but as I can not prove the optimal policy I can not test my assumption.

Used α : A small value of 0.1 for α is used. This small learning rate promises a slow but good convergence. If α is chosen to big the convergence is not guaranteed. In contrast to this a small α needs much more steps because the updates are very small and so the convergence is very slow. Therefore 5000 iterations steps are used in *WalkQLearning*.

Outcomes of changing ϵ or using pure greedy policy: If one uses a too high value or too low value of ϵ it may be that the robot only walks with one leg or does other things which are not desired. It may randomly find a better policy than with mostly greedy steps but in general my tries have shown that values over 0.4 lead to a worse behavior. A pure greedy policy leads in




Figure 6. Example for a pure greedy behavior

many cases to a loop of actions and states. One could improve that by choosing another reward matrix and lead the robot to the right cycle, but this would not be learning any more. A pure random policy may also generate a loop, but also actions which are not desired. The outcomes are not always the same for the same start state. Behaviors as mentioned above like the walking with only one leg may occur but also many times a good behavior is learned. Therefore it is important to choose




Figure 7. Example for a pure random policy

a proper ϵ . It must be small enough to find the actions which will get the most reward to learn the walking. It also must allow random steps so that the algorithm is

not stuck in local convergence.

Steps needed: I only can assume the number of steps which are necessary to find the optimal policy as i am not able to classify the optimal policy. I used 5000 iterations to achieve good results. 500 steps are to few to find a good policy with the small learning rate α which was used. I also did not see any improvement in increasing the steps over 5000.

Results of *WalkQLearning(s)*:

As one starts with a random initialisation and a partly random ϵ -greedy policy the results are not always the same in every try. here are the best results the algorithm delivered. The plots were generated with $\alpha = 0.1$, $\gamma = 0.999$, $\epsilon = 0.1$, $Steps = 5000$



Figure 8. Results from starting state $s = 16$




Figure 9. Results from starting state $s = 5$




Figure 10. Results from starting state $s = 12$